

1 R Basics

Exercise 1 (Data structures and subsetting)

- (a) Create a `list` named `list1` with the following components

`x` : a numeric vector 35, 34, 33, ..., 7, 6, 5, 6, 7, ..., 33, 34, 35 (don't type the single numbers!)

`y` : a factor with the first five letters of the Latin alphabet as levels, each repeated 12 times, except the last one which should be repeated 13 times

`mat` : a numeric 4×4 matrix whose entries are randomly drawn from an exponential distribution with $\lambda = 5$. Use the `set.seed()` function to make your result reproducible.

`list2` : a list with the components

- `t` : a numeric variable with the value 35
- `d` : a data.frame with the variables
 - `gender` : a factor with elements male, male, male, female, female, female, male, male, male, female, female, female
 - `age` : a numeric vector with elements 23, 48, 37, 37, 19, 54, 21, 20, 41, 26, 35, 32

- (b) Use subsetting operators to access the following information

- (a) the 4th and the 7th element of the vector `x`
- (b) the entry on position (3, 4) of the matrix `mat`
- (c) the first six elements of the first column of the data frame `d`
- (d) the age of the female individuals

What is the difference between the usage of `[]`, `[[]]` and `$` ?

- (c) Modify the created objects in the following ways

- (a) redefine the levels of the factor `y` from lowest to highest as 'c', 'd', 'b', 'a', 'e'
- (b) eliminate the first row and the third column of `mat`
- (c) add a column `age2` to `d` with a factor taking the value 'old', if the individual's age is above the threshold `t`, and the value 'young' otherwise
- (d) exclude the individuals that are at over (or exactly) 50 and strictly younger than 21 from the study

Exercise 2 (Loops and data manipulation)

- (a) Create a matrix X of dimension 2500×12 containing random numbers $x_{ij} \sim \mathcal{N}(\mu = 0, \sigma^2 = 3)$.
- (b) Create a new matrix $(y_{ij})_{ij}$ based on the one defined above, in which each column is standardized, i.e. $y_{ij} = \frac{x_{ij} - \bar{x}_j}{\sigma_j}$, where \bar{x}_j is the sample mean of the j -th column, and σ_j its sample standard deviation.

- (c) From X keep only the rows for which at least 6 out of the 12 values are greater than 0. Hint: Remember that a logical vector can be converted to a numeric vector where `FALSE` takes the value 0, and `TRUE` the value 1.
- (d) Check, if the elements of the 3rd column of `mat` from Exercise 1 are within the range defined by the elements in the first two columns. The output should be a logical vector.
- (e) In each row of X replace the maximum value with the minimum value.
- (f) Write a for-loop to do the following calculation for $i = 1, \dots, 2500$:
- if the i -th element of the first column X is positive, then add to the subsequent element $x_{i+1,j}$ a random number $u \sim \mathcal{U}(-1, 1)$
 - otherwise, if x_{ij} is negative, add to the subsequent element $x_{i+1,j}$ a random number $u \sim \mathcal{U}(-2, 2)$.
- (g) Consider the data.frame resulting from

```
set.seed(2015)
w <- runif(10)
x <- runif(10)
y <- runif(10)
z <- runif(10)
dat <- data.frame(a = w, b = x, c = y, d = z).
> dat
```

	a	b	c	d
1	0.06111892	0.70285557	0.6919100	0.75185674
2	0.83915986	0.39172125	0.4067718	0.25684487
3	0.29861322	0.03306277	0.2109400	0.38967137
4	0.03143242	0.40940319	0.6652073	0.88448520
5	0.13857171	0.74234713	0.7377556	0.57390551
6	0.35318471	0.88301877	0.9190050	0.18367673
7	0.49995552	0.26623321	0.8734601	0.11168811
8	0.07707116	0.07427093	0.8012774	0.32047459
9	0.65134483	0.81368426	0.5243978	0.09095567
10	0.51172371	0.38194719	0.1272213	0.47959896

Now imagine in each row the entries define two intervals $[a, b]$ and $[c, d]$. Add a column to `DF` with entry `TRUE`, if the intervals overlap and `FALSE`, if they don't.

Find ways to avoid for-loops in (a)–(c). Useful functions might be

`rnorm()`, `runif()`, `apply()`, `mean()`, `sd()`, `scale()`, `rowSums()`, `min()`, `max()`.

If you are not familiar with some of these functions, use `?function` to check the documentation. Compare the `system.time()` of the solutions with and without for-loops.

Exercise 3 (Basic graphical tools)

- (a) Consider again the matrix X from Exercise 2 (a) and create the following plots:
- A scaled histogram of the first column. Also add a dashed red line representing the estimated density.
 - A quantile-quantile plot comparing the sample quantiles of the first and second column. Also add a line with intercept 0 and slope 1 to improve comparability. Try different `pch`-values and colours.
 - Boxplots of the first, fourth and seventh row in one plotting window. Rename the group labels to `blue`, `green` and `yellow`.

2 Advanced Graphics

Exercise 1 (From basic plots to complex graphics)

Read the data set `school_math.raw`. It describes the mathematical achievements (`mathach`) of male and female (`Sex`) students at 4 different schools (`school`). Additionally there is information about the socio-economic status (`ses`) of the students family and the sector (`Sector`) of the school (public or catholic).

- (a) Get familiar with the data set, e.g. via the functions `summary()` or `head()`.
- (b) Create the following four plots in only one plotting window:
- a **scatterplot** of the achievement vs. the socio-economic status. Add a line representing the estimated linear relationship between those variables.

a **histogram** of `mathach`.

a **scaled histogram** of `ses`. Add a line with the estimated density.

boxplots of the mathematical achievement for each school separately.

First, don't modify any options. In a second step include the following variations:

- The color of the numbers on the axes should be red.
- Data points should be represented by filled squares (■) instead of circles (○).
- Use squared plotting regions for each plot.
- All occurring lines should be dashed.

Modify your plot by changing the settings within the single plot functions and in the overall graphics options using `par()`.

Hint: Make a backup of the default `par()` settings, such that you can reset the settings to default afterwards.

Hint: The documentations of `plot`, `points`, `lines`, `abline`, `par`, `hist`, `boxplot` might be helpful.

- (c) Reset the graphic settings to the default values.

Exercise 2 (The layout function)

Add additional information to the scatterplot from Exercise 1 in form of boxplots of the data at the axes. For this purpose, get familiar with the function `layout()` and its options. Also use `layout.show()` to visualize different settings.

Exercise 3 (Lattice plots and grouped data)

- (a) Estimate regression lines of the form

$$\text{mathach} = \beta_0 + \beta_1 \cdot \text{ses} + \epsilon$$

for the different schools separately. Plot your point estimates against the school. Use one plotting window for the intercepts and another one for the slopes. The y -axes should be labeled β_0 and β_1 , respectively.

Hint: Look at the function `lmList()` from the package `nlme`. Use `expression()` for the labeling.

- (b) Plot the mathematical achievement (in dependence of the `ses`) of boys as green triangles and of girls as red squares. Add a legend to the plot indicating this grouping.
- (c) Load the packages `lattice` and `nlme`. For the school data set, we want to compare the data in the different groups. For that purpose create a `groupedData` object with response `mathach`, covariate `ses` and school as grouping variable. Now

- Plot the new object. What is the difference between the results from `plot(data)` and `plot(data_new)`, where `data` is the original data set without grouping structure and `data_new` is the `groupedData` object?
- Create box-whiskers plots (`bwplot()`) of the residuals from an overall linear model according to `mathach ~ ses` grouped by the different schools.
- Introduce a random intercept for each school and estimate a mixed model using the function `lme()`. Compare the results with your findings from (a). Try `compareFits()` and `comparePred()`.
- Plot the confidence intervals of the estimated intercepts and slopes for the different schools.

3 Data Management

Exercise 1 (Spreadsheet-like data)

- Load the file `knee.txt` using the command `read.table()`. Compare what happens when using `header = TRUE` and `header = FALSE`.
- What class do the single variables have in R? Are the assigned classes reasonable in all cases? Define the classes of `th`, `gen` and `pain` as factor directly when reading in the data.
- Additionally rename the variables to *treatment*, *age*, *sex* and *agony*.
- Open the data set `knee2.txt` first in a text editor. What do you observe? Load the data into R such that missing values are automatically identified as those.
- Do the same with `knee3.txt`. What is the difference in contrast to `knee2.txt`?
- Add a column to the data frame `knee2` with the value *old*, if `age` is larger than 40 and *young* otherwise. Save the resulting data set in the following formats:
 - `*.Rdata`,
 - `*.raw` without quotes and with `:` as a separator,
 - `*.dat` without column or row names,
 - `*.csv`.

For the latter compare the results from `write.table`, `write.csv` and `write.csv2`.

Exercise 2 (Handling date objects and ordering data in R)

The dataset `movies.csv` contains information on the 10 most successful movies of each for the years 2012-2015.

- What is the class of the variables `release` and `end`? Obviously those two variables are the dates of the release and the last day of the movies in cinemas. Use `as.Date()` to transform them into dates.
- Add a column to the data frame with the number of days the movies have been shown in cinemas. Add another column with the number of complete weeks.
- Add a column with the year of release.
- Rank the movies according to the number of weeks the movies have been shown in cinemas.

- (e) (optional) Write a piece of code which has as output the name of each year's most successful movie (according to its gross income).
- (f) (optional) Find out on which `weekday()` Carl Friedrich Gauss was born.

Exercise 3 (Converting objects and *.RData-objects)

Load the data stored in `data.Rdata`. Use `ls()` to find out which objects are stored in that file. Convert the list into a vector and the data frame into a matrix. What do you observe? Try to explain what happens.

Exercise 4 (Foreign data structures)

- (a) The file `blutdruck.sav` is a SPSS data file. Use the according function from the package `foreign` to load it into R. What do you need to do to directly create a data frame in R?
- (b) Use the function `read.sas7bdat()` from the package `sas7bdat` to read in the SAS file `s05_01.sas7bdat`.
- (c) Read in the Stata file `golf.dta` with `read.dta()`.

4 Functions, Debugging & Condition Handling

Exercise 1 (Simple functions and restrictions)

In this exercise, we want to build a function to calculate the binomial coefficient of two integers n, k

$$\binom{n}{k} = \frac{n!}{k!(n-k)!} \quad (1)$$

Here, we will treat the binomial coefficient as not defined for negative numbers and non-integers¹. Follow the next steps carefully:

- (a) What are the definition and basic properties of the factorial of an integer? Also recall the properties of (1) for two integers.
- (b) Write a simple function `my_factorial()` which calculates the factorial of an integer n . Don't add any restrictions on n yet. Also, evaluate `my_factorial` at numbers from $\mathbb{R} \setminus \mathbb{N}$. What do you observe?
- (c) Let your function return an error message, if n is not an integer. The built-in function `is.integer()` will not help you in this case (why?). Hence, write your own function which returns `TRUE` for integers and `FALSE` for non-integers. Check if your function fulfills the properties from (a). If not, add necessary exceptions.
- (d) Write a function `my_binomial()` to calculate the expression in (1). Make use of your results from (b) and (c). Evaluate the expression `my_binomial(2, 3)`. Is the resulting error message reasonable? Use `traceback()` to locate where the error occurred.
- (e) Finally, add an exception for the case $k > n$ to `my_binomial()`. The function should then return the value 0. All other restrictions should be kept as before.

Exercise 2 (Scope and environments)

For a given vector $x = (x_1, \dots, x_n)$, the function

¹Note that there are extensions based on the Γ -function for these cases.

```
running_mean = function(k) {
  1 / k * sum(x[1:k])
}
```

calculates the running mean of order $k \leq n$, i.e.

$$\bar{x}_k = \frac{1}{k} \sum_{j=1}^k x_j$$

Guess the output of the following two scenarios. Then, check if your prediction was correct.

- (a) `x = 1:5`

```
running_mean = function(k) {
  1 / k * sum(x[1:k])
}
running_mean(2)
```
- (b) `x = 1:5`

```
rm(list = ls())
running_mean = function(k) {
  1 / k * sum(x[1:k])
}
running_mean(2)
```
- (c) Explain what went wrong in (b). Modify the function `running_mean` to make it independent from the environment/workspace.
- (d) Add reasonable restrictions on x and k .

Exercise 3 (Closures and environments)

In contrast to problems caused by different environments as in the previous exercise, one can explicitly use these environments.

1. Write a function `n.root` that allows for generating the family of n -th root functions

$$\sqrt[n]{x} = x^{\frac{1}{n}}$$

For that purpose write a function within a function, such that x is used in the inner environment and n in the outer environment.

2. Can you use `n.root` to explicitly calculate the roots of a number?
3. Based on `n.root` write functions which calculate the square and cubic root, respectively.
4. Use `n.root` to calculate the first 10 roots of $x = 500$.