



Applied Statistics: R

Semester 2018-I

Francisco Rosales Marticorena

Faculty of Economics & Finance

Table of Contents:

1 Introduction to R

2 Advanced Graphics

Introduction to R

Week 1:

- 1 R basics, GitHub and \LaTeX
- 2 graphical methods in R

Week 2:

- 1 data management (import/export)
- 2 functions, debugging, condition handling

Week 3:

- 1 profiling, performance and parallelization
- 2 numeric methods and simulations

If possible there is time: Rcpp, building packages.

Classes:

- One exam on mid-term week. 20% of your grade.
- More an R course than an “applied stats” course.
- 3 weeks, 5 hours \times week.
- TA comes the 3rd hour of the three-hour block.

Coordination:

- No Blackboard.
- Course repository [here](#).
- Use the Homework file to upload solutions of exercises.

Reading Material

- Hadley Wickham. **Advanced R**. The R series. CRC Press, 2015.
Available online: <http://adv-r.had.co.nz/>.
- Owen Jones, Robert Maillardet, and Andrew Robinson. **Introduction to scientific programming and simulation using R**. Chapman & Hall/CRC, 2009.
- Brian Everitt and Torsten Hothorn. **A handbook of statistical analyses using R**. Chapman & Hall/CRC, 2006.
- R Data Import/Export manual. Available online: <http://cran.r-project.org/doc/manuals/r-release/R-data.pdf>
- Deepayan Sarkar. **Lattice: multivariate data visualization with R**. Use R! Springer, 2008.
- Roger S Bivand, Edzer J Pebesma, and Virgilio Gómez-Rubio. **Applied spatial data analysis with R**, volume 10 of Use R! Springer, 2008.

Who are you?

- Ever used R?
- Text editor for R: Rstudio, emacs, Vim?
- Github, L^AT_EX?
- Operating system: Linux, Apple, Windows?
- Other programming languages: C, C++, Java, Python, Fortran, Julia, Matlab, Mathematica?

Heard of Laboratoria? kind of like that...

- Worldwide organization.
- Empowers women through coding.
- Third meet up on April 27th at UP Aula Magna H 15:00.
- Topic: Text-mining in Finance with R.
- Speaker: Leda Basombrío, Strategic Analysis Manager at BCP.

Not mandatory, but recommended. Not only for women.

help:	<code>?topic, help(topic), args(some function)</code>
assignments:	<code>x <- 5, x = 5, 5 -> x</code>
operators:	<code>+, -, *, /, ^, &, &&, , </code> ; see <code>help("+")</code>
comparisons:	<code>==, !=, >, >=, <, <=</code> ; see <code>help("==")</code>
loops:	<code>for, while, repeat</code>
comments:	everything that follows <code>#</code>

case sensitive: usage of CAPITAL and small letters matters!

Basic Data Structures

numeric (\mathbb{R}):	1, 301L, .141, 1.23e-3, NaN, Inf, -Inf
complex ($\mathbb{C} \setminus \mathbb{R}$)	1+0i, 1i, 3+5i
logical:	TRUE, FALSE, NA
character:	"hello", "I'm a string"
numeric:	no distinction between integers and doubles
missing:	stored as NA, and are logical.

Can check the basic structure with `str()`, `mode()` or `storage.mode()`.

```
str(1)          ## num 1
mode(1)         ## [1] "numeric"
storage.mode(1) ## [1] "double"
storage.mode(1L) ## [1] "integer"
mode(pi)        ## [1] "numeric"
storage.mode(pi) ## [1] "double"
str(LETTERS)    ## chr [1:26] "A" "B" "C" "D" "E" ...
```

Construction & Coercion

- vector can be constructed with `c()`
- coerced to a type `xxx` by `as.xxx()`
- when combining different data types, they will be coerced to the most flexible type
- coercion often happens automatically
- check if `xxx` is a specific type by `is.xxx()`

```
storage.mode(c(1,2L))      ## [1] "double"
storage.mode(c("a",1))     ## [1] "character"
x <- c(FALSE, TRUE, NA)
as.numeric(x)              ## [1] 0 1 NA
# Total number of TRUEs
sum(x, na.rm = TRUE)       ## [1] NA
```

In R, there three ways to represent “nothing”, but the reason for the missingness of the information can be distinguished:

NA	missing sample values
NaN	wrong math, e.g. $\log(-1)$, $1/0$
NULL	null pointer

```
c(3, NA)      ## [1] 3 NA
c(3, 0/0)     ## [1] 3 NaN
c(3, NULL)    ## [1] 3
max(3, NA)    ## [1] NA
```

Some math operations can be performed with `Inf` and `-Inf`:

```
max(3, Inf)
```

```
## [1] Inf
```

```
min(3, Inf)
```

```
## [1] 3
```

```
c(Inf + Inf, (-Inf) * Inf, Inf - Inf)
```

```
## [1] Inf -Inf NaN
```

Convolved Data Structures

- data structures in R can be organized by their dimensionality and if all their contents are of the same type (or not):

object	dimension	homogeneous	heterogeneous
1d	length()	atomic vector	list
2d	dim()	matrix	data frame
nd	dim()	array	–

- a `data.frame` is a matrix with different data type columns.
- a `list` can have different data type elements.

```
x <- matrix(1, nrow = 5, ncol = 2)
is.matrix(x)           ## [1] TRUE
as.vector(x)           ## [1] 1 1 1 1 1 1 1 1 1 1
x <- list(a = "Hallo", b = 1:10, pi = pi)
```

All objects can have arbitrary additional attributes, used to store metadata about the object.

- can be thought of as a named list (with unique names); other frequently encountered attributes: “dimnames”, “names”, “class”(!)
- can be accessed all at once (as a list) with `attributes()`, or individually with `attr()`.
- arrays are simply vectors with a “dim” attribute.
- factor is a vector with the “levels” attribute
- `as.xxx()` functions delete all attributes including dimensionality

Attributes

```
x <- matrix(1:10, ncol = 5)
attributes(x)                ## $dim
                             ## [1] 2 5

rownames(x) <- c("Eins", "Zwei")
attributes(x)                ## $dim
                             ## [1] 2 5
                             ## $dimnames
                             ## $dimnames[[1]]
                             ## [1] "Eins" "Zwei"
                             ## $dimnames[[2]]
                             ## NULL

as.character(x)              ## [1] "1"  "2"  "3"  "4"  "5"...
attributes(as.character(x))  ## NULL
```


- There are three subsetting operators: `[`, `[[`, and `$`
- the three types of subsetting:
 - Positive integers return elements at the specified positions.
 - Logical vectors select elements where the corresponding logical value is `TRUE`; application of logical expressions.
 - character vectors to return elements with matching names.
- important differences in behaviour of different objects (e.g., vectors, lists, factors, matrices, and data frames).
- More advanced subsetting, in particular in combination with convoluted logical expressions, can be done using the functions `subset()` and `which()`.
- The default `drop=TRUE` simplifies the data type of the result.

Atomic Vectors

Use “[”-operator and number, logical vector or name of the element you want to pull out.

```
x <- c(2.1, 4.2, 3.3, 5.4)
x[c(3, 1)]          ## [1] 3.3 2.1
x[-c(3, 1)]         ## [1] 4.2 5.4
x[c(TRUE, TRUE, FALSE, FALSE)] ## [1] 2.1 4.2

(y <- setNames(x, letters[1:4])) ##  a    b    c    d
## 2.1 4.2 3.3 5.4
y[c("d", "c", "a")]      ##  d    c    a
## 5.4 3.3 2.1
```

Subsetting matrices and arrays with “[”-operator like vectors, while the dimension is separated by comma:

```
x <- matrix(1:10, ncol=2)
colnames(x) <- c("Eins", "Zwei")
```

```
x[1:2,]           # results a matrix
x[,"Zwei"]         # results a vector
x[,"Zwei", drop = FALSE] # results a matrix
x[-3,]            # everything, but not third row
```

Subsetting lists with "["-operator returns always a list, while [[, and \$ pull out elements of the list:

```
x <- list(a = "Hallo", b = 1:10, pi = pi)
```

```
x$a      # first element of the list
```

```
x[['a']]
```

```
x[[1]]
```

```
x[1]     # list with one element
```

```
x[2:3]   # list with two elements
```

```
x[[2:3]] # wrong result
```

Data Frames

Data frames possess the characteristics of both lists and matrices: if you subset with a single vector, they behave like lists; if you subset with two vectors, they behave like matrices.

```
iris[1:10,]           # data frame with 10 rows
iris[,1]              # numerical
iris$Sepal.Length     # the name
iris$Sepal.Length     # Oops! what happened?
iris["Sepal.Length"]  # again first column
iris["Sepal.Length"]  # Error:  undefined columns selected
iris[,1, drop = FALSE] # data frame with one column
```

- conditional evaluation: if, else, ifelse
- loops: for, while, repeat, switch

basic vocabulary:

if, &&, ||

for, while, repeat

next, break

switch

ifelse

```
if (<test>) {  
  <expression1>  
} else {  
  <expression2>  
}
```

- else block is optional
- <test> has to result in a value that is TRUE or FALSE
- only the first element of <test> is used, ow a warning is triggered
- to eval more than one statement use &, | or all() and any()
- can be nested

for Loop

```
for (<var> in <vector>) {  
  <expression>  
}
```

```
sum <- 0  
for(i in 1 : length(x)) {  
  sum <- sum + x[i]  
}
```

```
sum <- 0  
for(x_value in x) {  ## more efficient  
  sum <- sum + x_value  
}
```

Use `seq_along(x)` instead of `1:length(x)`

while Loop

```
while(<test>) {  
  <expression>  
}
```

E.g., the sum until the first NA:

```
sum <- 0  
i <- 1  
while((i <= length(x)) && !is.na(x[i])) {  
  sum <- sum + x[i]  
  i <- i + 1  
}
```

Be aware of infinite loops!

- next jumps to the next iteration in for or while loops
- break terminates for or while loops.

```
x <- c(1, 1, 1, NA, 2)
sum <- 0
for(val in x) {
  if(is.na(val)) break
  sum <- sum + val
}
```

```
x <- c(1, 1, 1, NA, 2)
sum <- 0
for(val in x) {
  if(is.na(val)) next
  sum <- sum + val
}
```

Style Example: Bad

```
fWLM<-function(y,X_mat,w){T0<-t(X_mat)%*%diag(w)%*%X_mat
t<-system.time({t_1<-solve(T0)%*%t(X_mat)%*%(w*y);
t2<-X_mat%*%t_1})
return(list(beta=t_1,hat=t2,stddev=sqrt(sum(w*(t2-y)^2))/
(length(y)-ncol(X_mat)), wts=w,t=t[[3]]))}
```

Style Example: Good

```
fit_weighted_lm <- function(response, design, weights) {  
  n_obs <- length(response)  
  n_coef <- ncol(design)  
  time_start <- Sys.time()  
  wcrossprod_design <- crossprod(design * weights, design)  
  weighted_response <- weights * response  
  coef <- solve(wcrossprod_design, t(design) %*% weighted_response)  
  time <- Sys.time() - time_start  
  fitted <- design %*% coef  
  residuals <- response - fitted  
  weighted_rss <- sum(weights * residuals^2)  
  sd_resid <- sqrt(weighted_rss / (n_obs - n_coef))  
  list(coef      = coef,  
        fitted   = fitted,  
        sd_resid = sd_resid,  
        weights  = weights,  
        time     = time)  
}
```

- find meaningful file names; if files need to be run in sequence, prefix them with numbers.

0-download.R

1-parse.R

2-explore.R

- avoid uppercase
- use an underscore to separate words within a name
- use nouns for variable names and verbs for function names

- strive for names that are concise and meaningful (this is not easy!).

# Good	# Bad
day_one	first_day_of_the_month
day_1	dayone
	djm1

- avoid using names of existing functions and variables.

```
# Bad
T <- FALSE
c <- 10
t <- temporal_variable
mean <- function(x) sum(x)
```

Formatting

- strive to limit your code to 80 characters per line.
- when indenting your code, use two spaces. Never use tabs.
- place spaces around all infix operators (=, +, -, <-, etc.), before parentheses, and after comma (just like in regular English)

Good

```
average <- mean(feet / 12 + inches, na.rm = TRUE)
```

Bad

```
average<-mean(feet/12+inches,na.rm=TRUE)
```

Good

```
if (debug) do(x)
```

```
plot(x, y)
```

Bad

```
if(debug)do(x)
```

```
plot (x, y)
```

- an opening (or closing) curly brace should always be followed by a new line, unless it's followed by `else`.

```
if (y == 0) {  
  log(x)  
} else {  
  y^x  
}
```

- use commented lines of `-` and `=` to break up your file into chunks.

```
# Load data -----  
# Plot data -----
```

- `formatR::tidy_source(source="input.R",file="output.R")`
cleans up and does some automatic formatting

Advanced Graphics

- creating interest and attention of the reader
- essential meaning can be visualized at a glance
- comprehensive picture of a problem gives more complete and balanced understanding
- the human visual system is very powerful in detecting patterns: outliers, diagnose models, search for perhaps unexpected phenomena

- A graphical device can be thought as a paper on which you can draw with different pens and colours, but nothing can be deleted.
- It can be opened more than one device, but there is only one active.
- There is no difference no matter which device is used.
- Typical steps to produce a graphic is:
 - 1 start device, e.g. `pdf('testgraphics.pdf')`
 - 2 generate graphic, e.g. `plot(1:10)`
 - 3 close device: `dev.off()`
- If no device is open, using a high-level graphics function will cause a device to be opened.

The following graphics devices are currently available:

- `pdf()`: write PDF graphics commands to a file; can be handy for distribution to cooperation partners, integration in PDF \LaTeX , or viewing many graphics
- `postscript()`: writes PostScript graphics commands to a file
- `bitmap()`: bitmap pseudo-device via ‘Ghostscript’ (if available).

Interactive plotting with GUI:

- `x11()`: The graphics device for the X11 windowing system
- `png()`: compressed Bitmap, without loss
- `jpeg()`: compressed Bitmap with information loss, optimized for pictures with many color shades

For more info see `?Devices`.

High-level Plots

High-level functions generate/initialize a graphic, e.g.:

<code>plot()</code>	depend of context
<code>barplot()</code>	Barplot
<code>boxplot()</code>	Boxplot
<code>coplot ()</code>	Conditioning plots
<code>contour()</code>	Contour line plot
<code>curve()</code>	Plotting functions
<code>dotchart()</code>	Dot Plots
<code>hist()</code>	Histogram
<code>image()</code>	Countour Plot (3. Dim. as color)
<code>mosaicplot()</code>	Mosaicplots (categorical data)
<code>pairs()</code>	Scatterplot matrix
<code>persp()</code>	perspective surface
<code>qqplot()</code>	QQ-Plot

High-level Plots

Many functions can be applied to different object types. They react in a “intelligent” way, so that a meaningful graphic can be found.

```
plot(trees)                                ## scatterplot matrix
plot(Volume ~ Girth, data = trees)         ## scatterplot

tree.lm <- lm(Volume ~ Girth, data = trees)
abline(tree.lm)                            ## regression line
plot(tree.lm)                              ## residual/diagnostic plots
boxplot(trees)                             ## boxplots
qqnorm(trees$Volume)                       ## quantile plot
```

- for nicer graphics, the `par` graphical parameters can be adapted
- some graphical parameters can be adjusted in high-level functions
- for help check `?plot`, `?plot.default`, or more comprehensive `?par`

Some graphical parameters can be set in high-level functions like `plot()`. For example:

<code>axes</code>	should the axes be plotted?
<code>col</code>	color
<code>log</code>	logarithmic scale
<code>main, sub</code>	title and subtitle
<code>pch</code>	symbol for points
<code>type</code>	type (l=line, p=point, b=both, n=none)
<code>xlab, ylab</code>	x-/y-axis label <code>xlim, ylim</code>
<code>xlim, ylim</code>	x-/y-axis range

Graphical Parameter

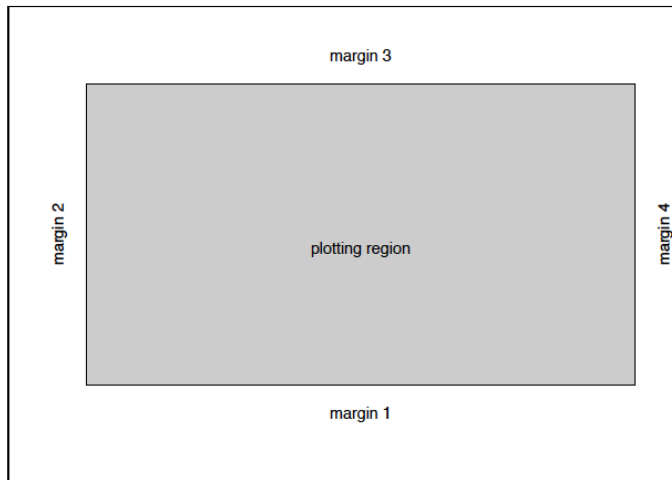
The most commonly used arguments in `par()`:

<code>bg</code>	background color
<code>cex</code>	size of a point or a letter
<code>las</code>	should labels be placed parallel wrt the axes
<code>lty, lwd</code>	line type (dashed, ...) and line width
<code>mar</code>	size of the margins
<code>mfc col, mfrow</code>	multiple plots in one device in rows/columns
<code>mfg</code>	which plot in a device should be chosen?
<code>oma</code>	size of the outer margins
<code>usr</code>	current extrema of the user coordinates
<code>xaxt, yaxt</code>	x-/y-axis scaling

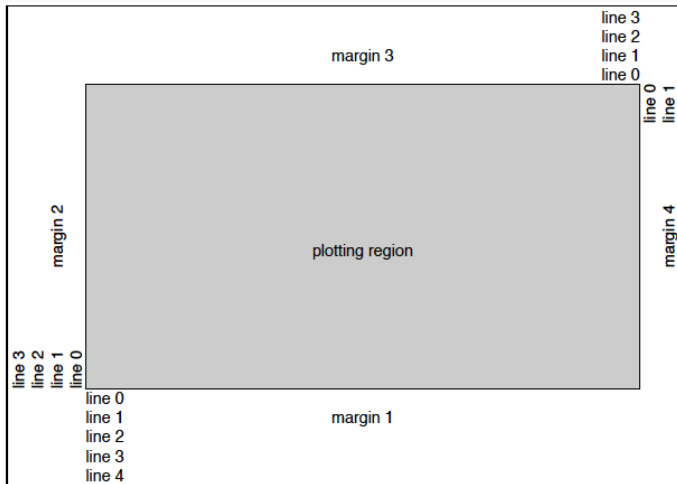
Graphical Parameter

```
opar <- par(mfrow = c(1, 1),bg = "White")
# example 1
par(mfrow = c(2, 2))
boxplot(trees, col = "blue")
hist(trees$Volume, las = 1)
qqnorm(trees$Volume, cex.axis = 2, pch = (trees$Girth > 14) + 8)
plot(trees$Girth,trees$Height,cex = scale(trees$Volume,center=FALSE))
par(opar)


# example 2
set.seed(123)
x <- rnorm(100)
par(bg = "lightgreen")
hist(x, freq = FALSE, col = "red", las = 1,
     xlim = c(-5, 5), ylim = c(0, 0.6),
     main = "100 Draws from N(0,1)-distributed random variables")
curve(dnorm, from = -5, to = 5, add = TRUE, col = "blue", lwd = 3)
par(opar)
```



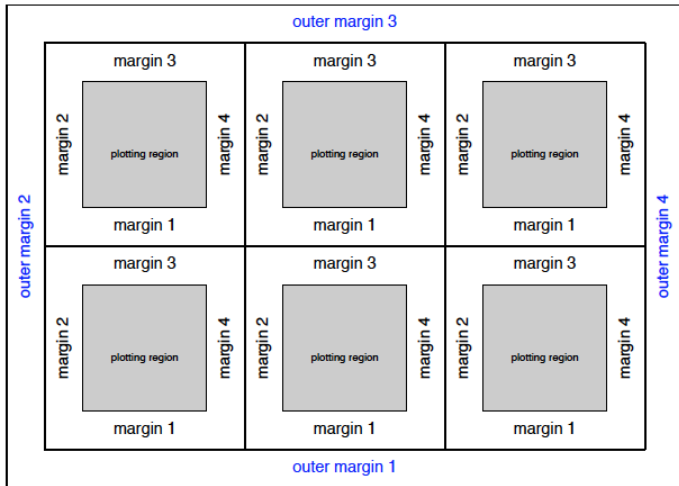
Device Control



Device Control



```
par(mfrow=c(2,3))
```



layout() Function

- layout() organizes independent plots on one plotting device, also in irregular grids
- boxes can have different widths
- neighboring boxes can be combined
- boxes can be left empty

```
m <- matrix(c(1,1,0,2), 2, 2)
```

```
m
```

```
##      [,1]  [,2]
```

```
## [1,]    1    0
```

```
## [2,]    1    2
```

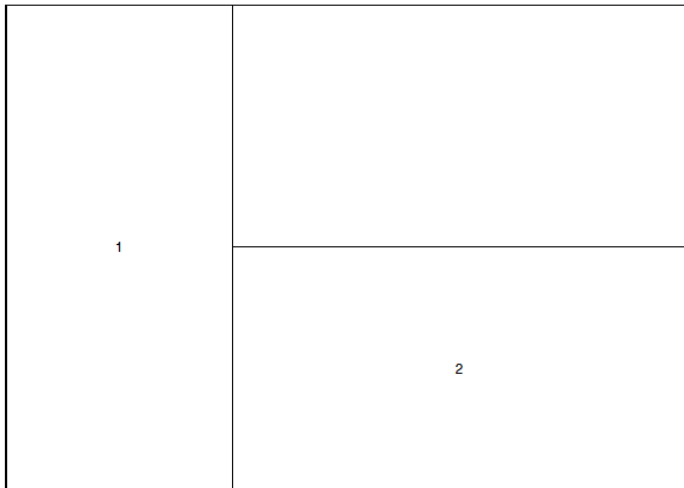
```
layout(m, widths=c(1,2))
```

```
x <- rnorm(100)
```

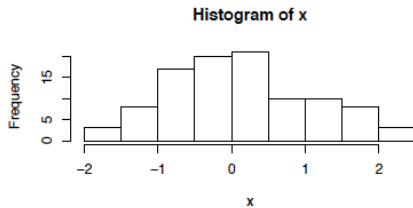
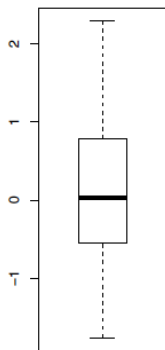
```
boxplot(x)
```

```
hist(x)
```

layout() Function



layout() Function



Low-level Graphics

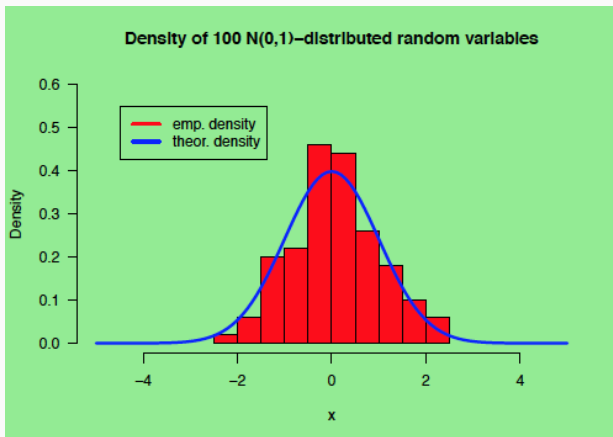
Low-level functions add elements to a (with high-level function) generated graphic, e.g., additional points, legends, etc.

<code>abline()</code>	“intelligent” lines
<code>arrows()</code>	arrows
<code>axis()</code>	axes
<code>grid()</code>	gridlines
<code>legend()</code>	legend
<code>lines()</code>	(stepwise) lines
<code>mtext()</code>	text in margins
<code>points()</code>	points
<code>polygon()</code>	(filled) polygons
<code>segments()</code>	vector lines
<code>text()</code>	text
<code>title()</code>	title label

Low-level Graphics

```
# example 2 (cont.):
```

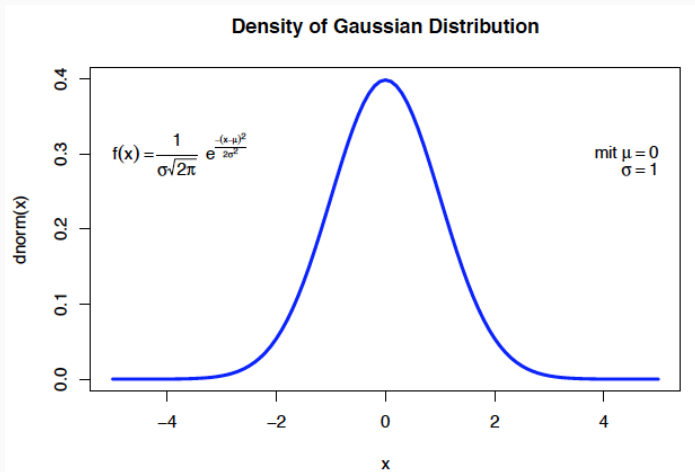
```
legend(-4.5, 0.55, legend = c("emp. density", "theor. density"),  
      col = c("red", "blue"), lwd = 3)
```



Mathematical Expressions

- mathematical notation and symbols formatted similar to \LaTeX code can be integrated in functions such as `axis()`, `legend()`, `mtext()`, `text()`, and `title()`
- For help check `?plotmath` or run `demo(plotmath)`

```
curve(dnorm, main = "Density of Gaussian Distribution",  
      from = -5, to = 5, col = "blue", lwd = 3)  
text(-3, 0.3,  
      expression(f(x) == frac(1, sigma * sqrt(2*pi)) ~  
                  e^{frac(-(x - mu)^2, 2 * sigma^2)}))  
text(4, 0.3, expression(paste("mit ", mu == 0)))  
sigma <- 1  
text(4.2, 0.28, bquote(sigma == .(sigma)))
```

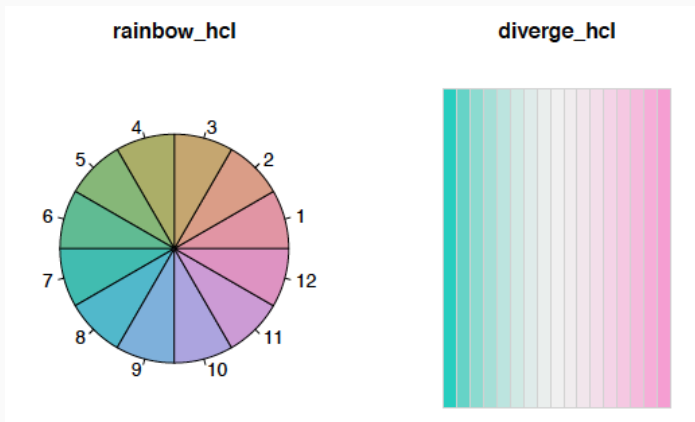


colorspace

The package `colorspace` provides various functions for perceptually-balanced color palettes

```
rainbow_hcl(12)
```

```
diverge_hcl(17, h = c(180, 330), c = 59, l = c(75, 95))
```

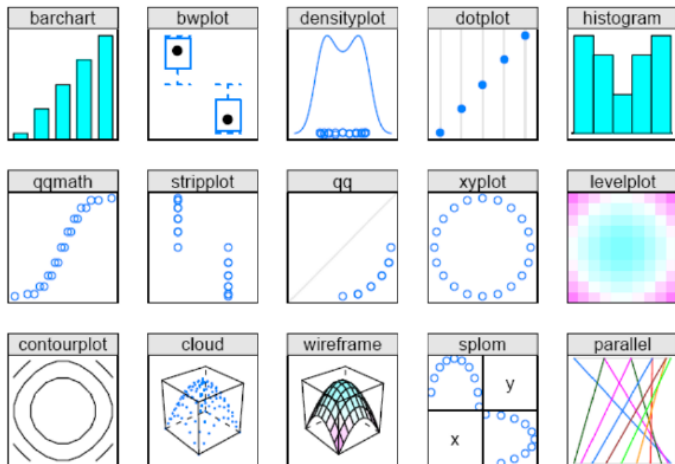


- trellis graphics is a family of techniques for viewing complicated data sets, that are based on basic concepts of human perception
- everything is possible (using a sufficient number of parameters)
- the trellis graphics system is in the `lattice` package
- the typical format is

```
graph_type(formula, data = )
```

<code>barchart()</code>	barplot
<code>bwplot()</code>	boxplot
<code>cloud()</code>	3D point clouds
<code>contourplot</code>	3D contour plot
<code>densityplot()</code>	kernel density plot
<code>dotplot()</code>	point plots
<code>histogram()</code>	histogram
<code>levelplot()</code>	levelplots
<code>panel.....()</code>	functions to add elements
<code>piechart()</code>	pie diagram
<code>print.trellis()</code>	plotting trellis object
<code>qq()</code>	QQ-plots
<code>stripplot</code>	strip plots
<code>wireframe()</code>	persp. 3D areas
<code>xyplot()</code>	scatterplot

Trellis Graphics



```
require(lattice)
attach(mtcars)
# create factors with value labels
gear.f <- factor(gear, levels = c(3,4,5),
  labels = c("3gears", "4gears", "5gears"))
cyl.f <- factor(cyl, levels = c(4,6,8),
  labels = c("4cyl", "6cyl", "8cyl"))

# kernel density plot
densityplot(~mpg,
  main = "Density Plot",
  xlab = "Miles per Gallon")
# kernel density plots by factor level
densityplot(~mpg|cyl.f,
  main = "Density Plot by Number of Cylinders",
  xlab = "Miles per Gallon")
```

The basic R system does not allow many possibilities for interactive graphics. Some exceptions are:

- `identify()` identifies a selected data point, e.g.:

```
x <- rnorm(10)
plot(x)
identify(x)
```

- `locator()` returns the coordinates of a selected point, which can be used for instance for the interactive placing of labels:

```
plot(x)
legend(locator(1), legend = "A legend", pch = 1)
```

Further References

- First analysis: Brian Everitt and Torsten Hothorn. [A handbook of statistical analyses using R](#). Chapman & Hall/CRC, 2006
- Trellis: Deepayan Sarkar. [Lattice: multivariate data visualization with R](#). Use R! Springer, 2008
- Colorspace: Achim Zeileis, Kurt Hornik, and Paul Murrell. [Escaping RGBland: Selecting colors for statistical graphics](#). Computational Statistics & Data Analysis, 53:3259-3270, 2009. doi: 10.1016/j.csda.2008.11.033